

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as py
import seaborn as sns
```

```
In [2]: d=pd.read_csv(r"C:\Users\user\Downloads\6_Salesworkload1 - 6_Salesworkload1.csv")
d
```

```
Out[2]:
```

	MonthYear	Time index	Country	StoreID	City	Dept_ID	Dept. Name	HoursOwn	HoursLease
<b>0</b>	10.2016	1.0	United Kingdom	88253.0	London (I)	1.0	Dry	3184.764	0.0
<b>1</b>	10.2016	1.0	United Kingdom	88253.0	London (I)	2.0	Frozen	1582.941	0.0
<b>2</b>	10.2016	1.0	United Kingdom	88253.0	London (I)	3.0	other	47.205	0.0
<b>3</b>	10.2016	1.0	United Kingdom	88253.0	London (I)	4.0	Fish	1623.852	0.0
<b>4</b>	10.2016	1.0	United Kingdom	88253.0	London (I)	5.0	Fruits & Vegetables	1759.173	0.0
...	...	...	...	...	...	...	...	...	...
<b>7653</b>	6.2017	9.0	Sweden	29650.0	Gothenburg	12.0	Checkout	6322.323	0.0
<b>7654</b>	6.2017	9.0	Sweden	29650.0	Gothenburg	16.0	Customer Services	4270.479	0.0
<b>7655</b>	6.2017	9.0	Sweden	29650.0	Gothenburg	11.0	Delivery	0	0.0
<b>7656</b>	6.2017	9.0	Sweden	29650.0	Gothenburg	17.0	others	2224.929	0.0
<b>7657</b>	6.2017	9.0	Sweden	29650.0	Gothenburg	18.0	all	39652.2	0.0

7658 rows × 14 columns



```
In [3]: d.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7658 entries, 0 to 7657
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   MonthYear       7658 non-null   object
1   Time index      7650 non-null   float64
2   Country         7650 non-null   object
3   StoreID         7650 non-null   float64
4   City            7650 non-null   object
5   Dept_ID         7650 non-null   float64
6   Dept. Name      7650 non-null   object
7   HoursOwn        7650 non-null   object
8   HoursLease      7650 non-null   float64
```

```

9 Sales units      7650 non-null float64
10 Turnover        7650 non-null float64
11 Customer        0 non-null   float64
12 Area (m2)       7650 non-null object
13 Opening hours   7650 non-null object
dtypes: float64(7), object(7)
memory usage: 837.7+ KB

```

In [4]: `d.head()`

Out[4]:

	MonthYear	Time index	Country	StoreID	City	Dept_ID	Dept. Name	HoursOwn	HoursLease	Sales units
0	10.2016	1.0	United Kingdom	88253.0	London (I)	1.0	Dry	3184.764	0.0	398560.0
1	10.2016	1.0	United Kingdom	88253.0	London (I)	2.0	Frozen	1582.941	0.0	82725.0
2	10.2016	1.0	United Kingdom	88253.0	London (I)	3.0	other	47.205	0.0	438400.0
3	10.2016	1.0	United Kingdom	88253.0	London (I)	4.0	Fish	1623.852	0.0	309425.0
4	10.2016	1.0	United Kingdom	88253.0	London (I)	5.0	Fruits & Vegetables	1759.173	0.0	165515.0

In [5]: `d.describe()`

Out[5]:

	Time index	StoreID	Dept_ID	HoursLease	Sales units	Turnover	Customer
count	7650.000000	7650.000000	7650.000000	7650.000000	7.650000e+03	7.650000e+03	0.0
mean	5.000000	61995.220000	9.470588	22.036078	1.076471e+06	3.721393e+06	NaN
std	2.582158	29924.581631	5.337429	133.299513	1.728113e+06	6.003380e+06	NaN
min	1.000000	12227.000000	1.000000	0.000000	0.000000e+00	0.000000e+00	NaN
25%	3.000000	29650.000000	5.000000	0.000000	5.457125e+04	2.726798e+05	NaN
50%	5.000000	75400.500000	9.000000	0.000000	2.932300e+05	9.319575e+05	NaN
75%	7.000000	87703.000000	14.000000	0.000000	9.175075e+05	3.264432e+06	NaN
max	9.000000	98422.000000	18.000000	3984.000000	1.124296e+07	4.271739e+07	NaN

In [6]: `d.columns`

Out[6]: Index(['MonthYear', 'Time index', 'Country', 'StoreID', 'City', 'Dept\_ID', 'Dept. Name', 'HoursOwn', 'HoursLease', 'Sales units', 'Turnover', 'Customer', 'Area (m2)', 'Opening hours'], dtype='object')

In [7]: `d.index`

Out[7]: RangeIndex(start=0, stop=7658, step=1)

```
In [8]: d=d.head(100)
d
```

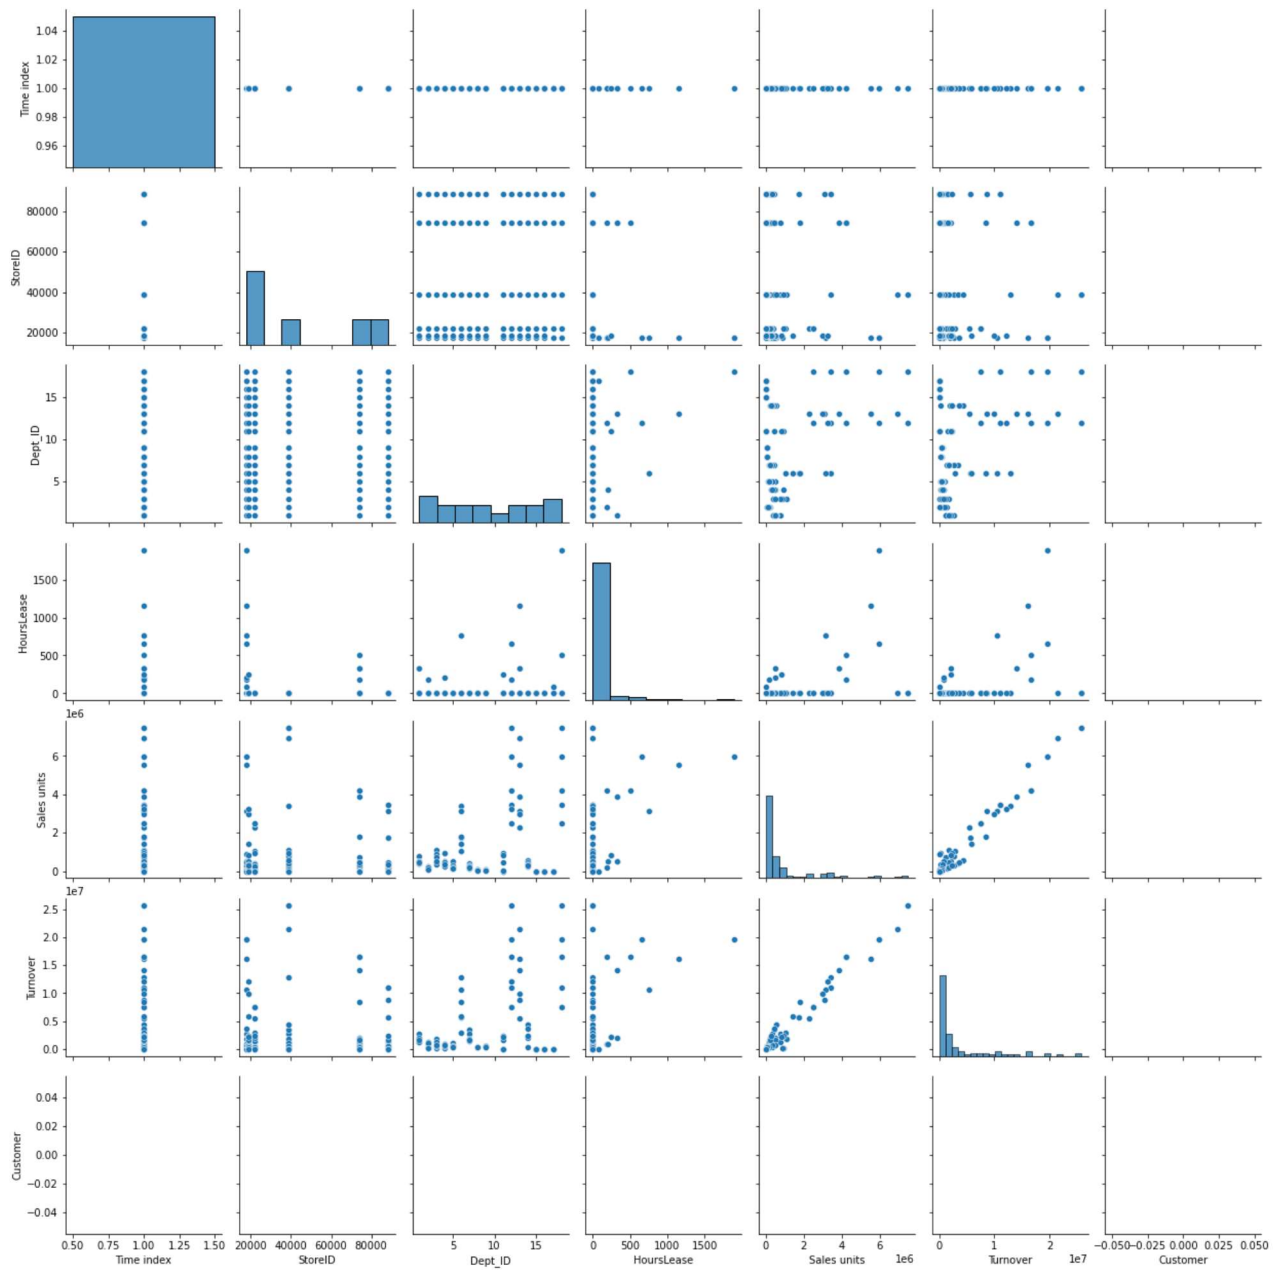
Out[8]:

	MonthYear	Time index	Country	StoreID	City	Dept_ID	Dept. Name	HoursOwn	HoursLease	Sal uni
0	10.2016	1.0	United Kingdom	88253.0	London (I)	1.0	Dry	3184.764	0.0	398560
1	10.2016	1.0	United Kingdom	88253.0	London (I)	2.0	Frozen	1582.941	0.0	82725
2	10.2016	1.0	United Kingdom	88253.0	London (I)	3.0	other	47.205	0.0	438400
3	10.2016	1.0	United Kingdom	88253.0	London (I)	4.0	Fish	1623.852	0.0	309425
4	10.2016	1.0	United Kingdom	88253.0	London (I)	5.0	Fruits & Vegetables	1759.173	0.0	165515
...	...	...	...	...	...	...	...	...	...	...
95	10.2016	1.0	United Kingdom	18808.0	London (II)	14.0	Non Food	7817.148	0.0	301500
96	10.2016	1.0	United Kingdom	18808.0	London (II)	15.0	Admin	5110.728	0.0	25
97	10.2016	1.0	United Kingdom	18808.0	London (II)	12.0	Checkout	6209.031	0.0	3262240
98	10.2016	1.0	United Kingdom	18808.0	London (II)	16.0	Customer Services	3115.53	0.0	25
99	10.2016	1.0	United Kingdom	18808.0	London (II)	11.0	Delivery	7209.777	246.0	843615

100 rows × 14 columns

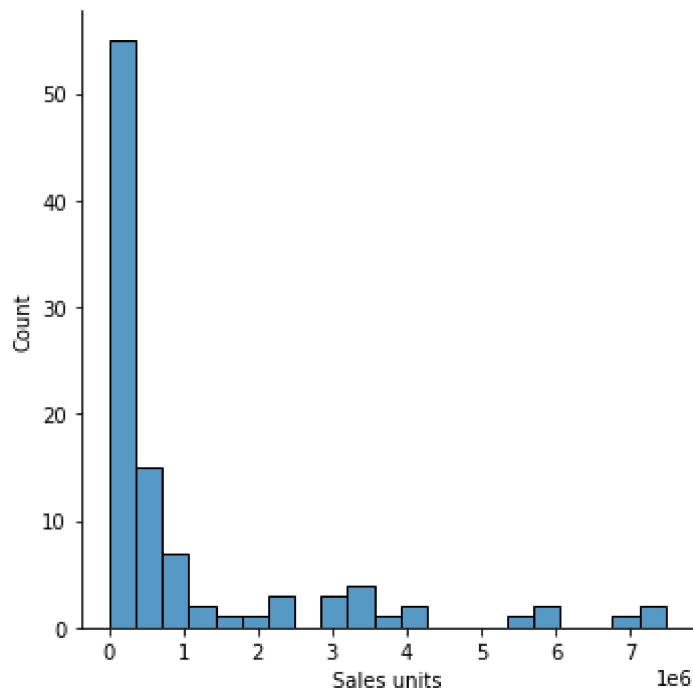
```
In [9]: sns.pairplot(d)
```

Out[9]: <seaborn.axisgrid.PairGrid at 0x22c1991e670>



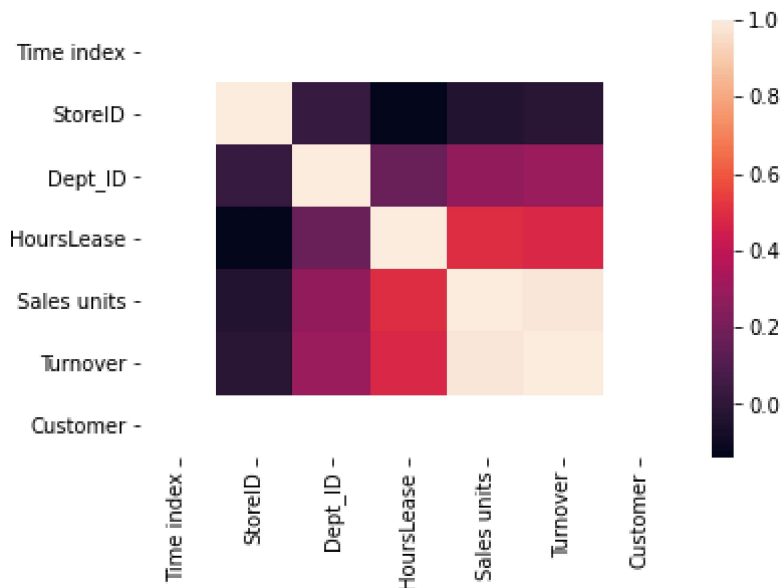
```
In [10]: sns.displot(d['Sales units'])
```

```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x22c1c099820>
```



```
In [11]: d1=d[['MonthYear', 'Time index', 'Country', 'StoreID', 'City', 'Dept_ID',
              'Dept. Name', 'HoursOwn', 'HoursLease', 'Sales units', 'Turnover',
              'Customer', 'Area (m2)', 'Opening hours']]
sns.heatmap(d1.corr())
```

Out[11]: <AxesSubplot:>



```
In [12]: x=d1[['MonthYear', 'Time index', 'StoreID', 'Dept_ID', 'HoursOwn', 'HoursLease', 'Turnover',
              'Sales units']]
y=d1['Sales units']
```

```
In [13]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [14]: from sklearn.linear_model import LinearRegression
```

```
In [15]: lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[15]: LinearRegression()
```

```
In [16]: print(lr.intercept_)
```

```
106150.93808447546
```

```
In [17]: coeff =pd.DataFrame(lr.coef_,x.columns,columns=["Co-efficient"])  
coeff
```


```
Out[17]:
```

	Co-efficient
--	--------------

<b>MonthYear</b>	0.000000e+00
<b>Time index</b>	-3.034548e-07
<b>StoreID</b>	-9.741773e-01
<b>Dept_ID</b>	-3.101410e+03
<b>HoursOwn</b>	-3.912454e-02
<b>HoursLease</b>	1.591050e+02
<b>Turnover</b>	2.984325e-01
<b>Area (m2)</b>	-5.137880e+00

```
In [18]: prediction =lr.predict(x_test)  
py.scatter(y_test,prediction)
```

```
Out[18]: <matplotlib.collections.PathCollection at 0x22c1e6b18b0>
```



```
In [19]: print(lr.score(x_test,y_test))
```

0.8560899084782179

```
In [20]: print(lr.score(x_train,y_train))
```

0.9764615700111092

```
In [21]: from sklearn.linear_model import Ridge,Lasso
```

```
In [22]: rr=Ridge(alpha=10)
         rr.fit(x_train,y_train)
```

Out[22]: Ridge(alpha=10)

```
In [23]: rr.score(x_test,y_test)
```

Out[23]: 0.8560817382584147

```
In [24]: la=Lasso(alpha=10)
         la.fit(x_train,y_train)
```

Out[24]: Lasso(alpha=10)

```
In [25]: la.score(x_test,y_test)
```

Out[25]: 0.8560897190216925

```
In [26]: from sklearn.linear_model import ElasticNet
         en=ElasticNet()
         en.fit(x_train,y_train)
```

Out[26]: ElasticNet()

```
In [27]: print(en.coef_)
```

```
[-0.00000000e+00  0.00000000e+00 -9.73228911e-01 -3.02266523e+03
 -5.49420387e-02  1.59182524e+02  2.98461871e-01 -5.15636356e+00]
```

```
In [28]: print(en.intercept_)
```

105535.9279539762

```
In [29]: print(en.score(x_test,y_test))
```

0.8560615377349597

```
In [30]: from sklearn import metrics
```

```
In [33]: print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,prediction))
```

Mean Absolute Error: 135639.1438373063

```
In [34]: print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

Mean Squared Error: 52132479159.11114

```
In [35]: print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

Root Mean Squared Error: 228325.38001525617

```
In [ ]:
```